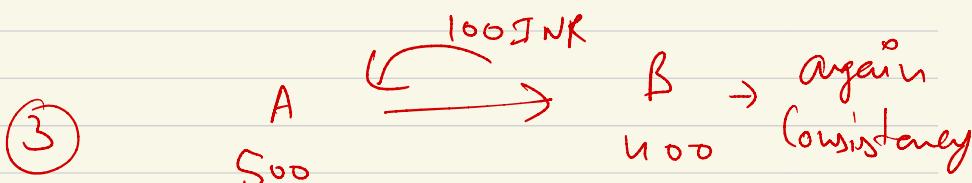
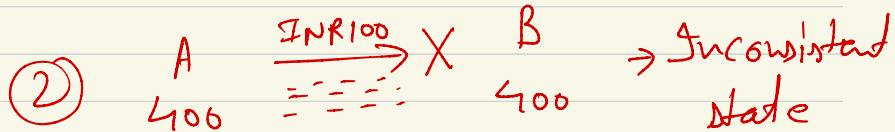
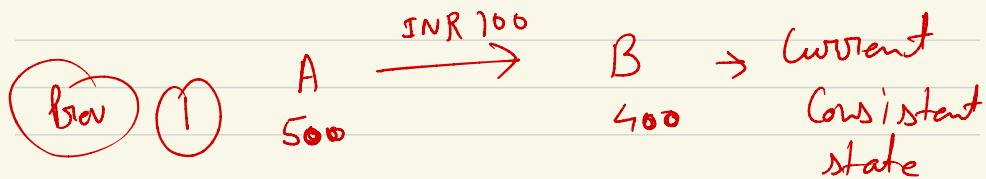
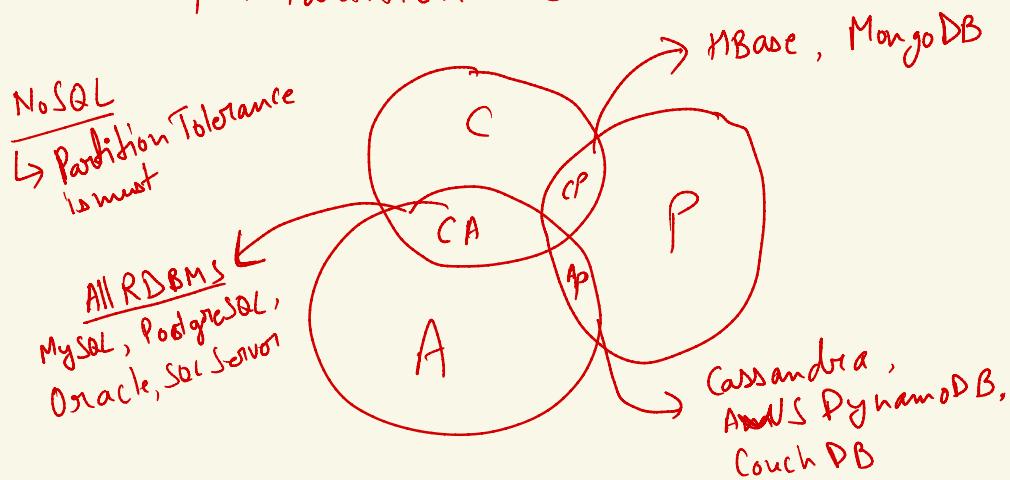


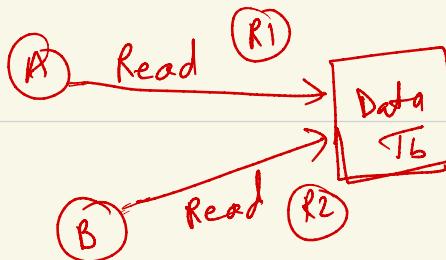

CAP Theorem → We can only achieve two properties from a database.

C → Consistency

A → Availability

P → Partition Tolerance



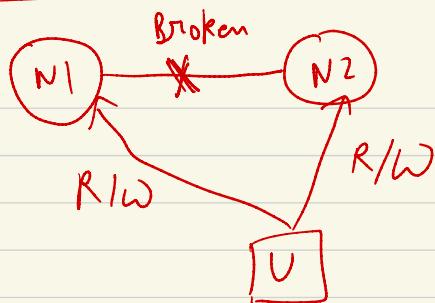


Availability

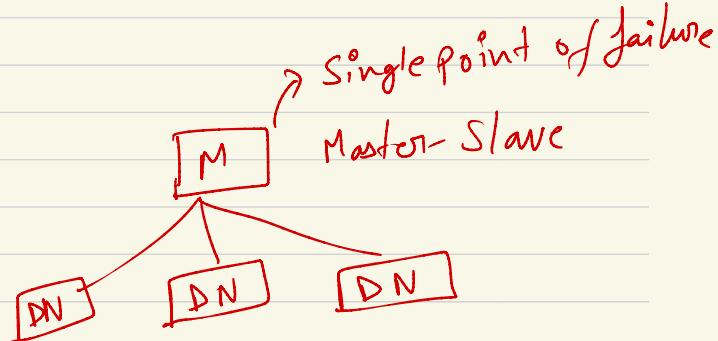
→ Our system should be up 24x7
 ↳ If we have single point of failure then it's hard to achieve

Partition-tolerance

If there is network break in the system still it should operate and serve request.



HBase

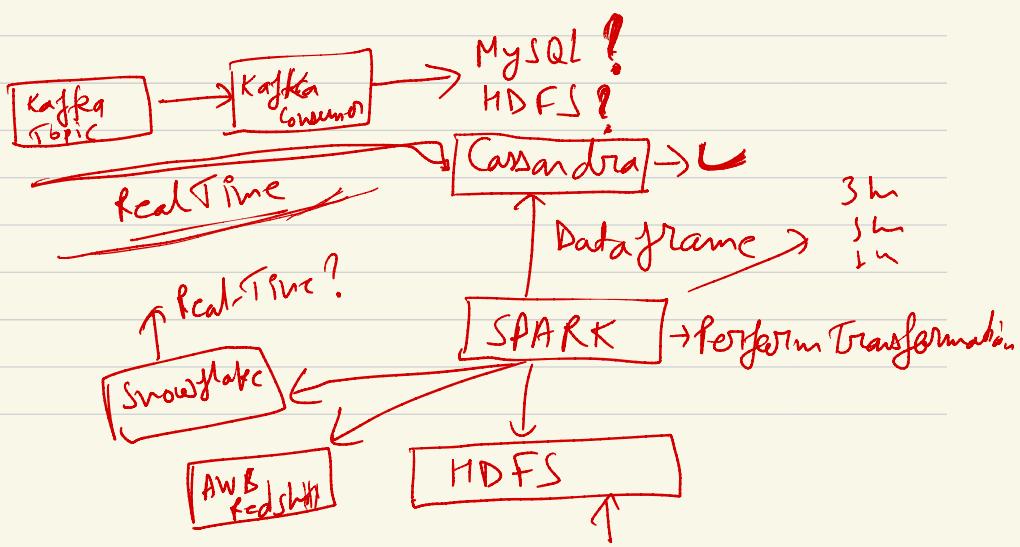


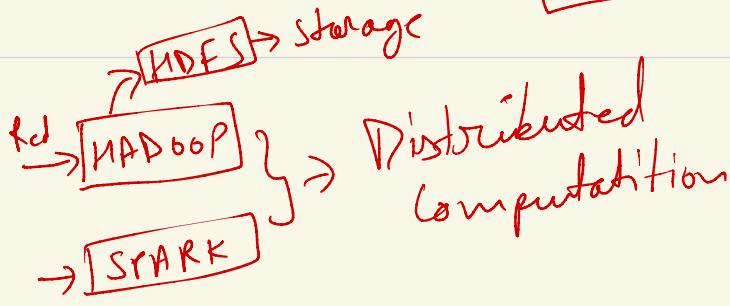
Cassandra

T

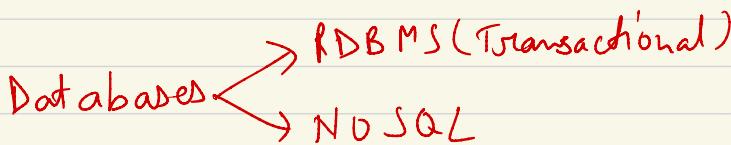
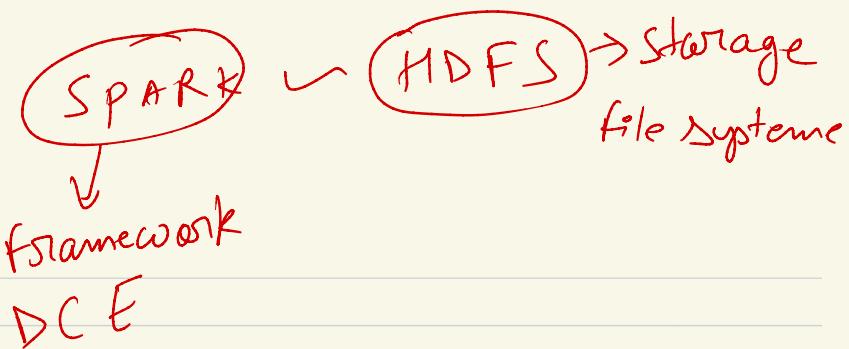
- It was developed by facebook
- It came as a open source project under Apache
- 3.0+, 3.25 Version
- Cassandra is not made for Analytical queries
- we can not perform joins, aggregation etc.
- It is beneficial for faster write operation
- Also beneficial for faster read search type queries
- Cassandra follow Ring Kind of Arch.

Real Time Data Pipeline (High Velocity)





C++, Java, Python, Scala



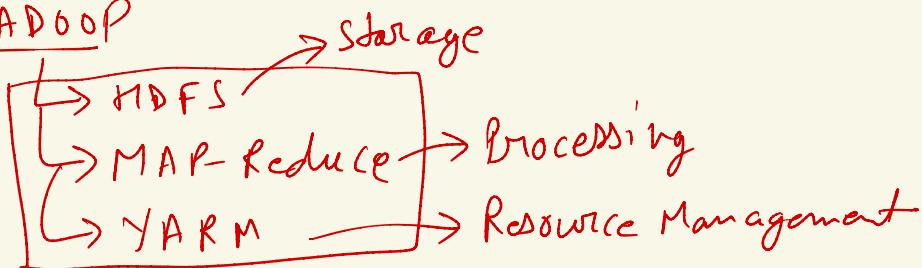
Messaging - Queue → KAFKA

Data Warehousing → Means storing transformed structured data
 ↳ Hive, AWS Redshift, Snowflake, Azure Data Factory

Distributed Computation Engine

- ↳ Process data in distributed fashion or parallel processing
- ↳ HADOOP
- ↳ SPARK
- ↳ Flink

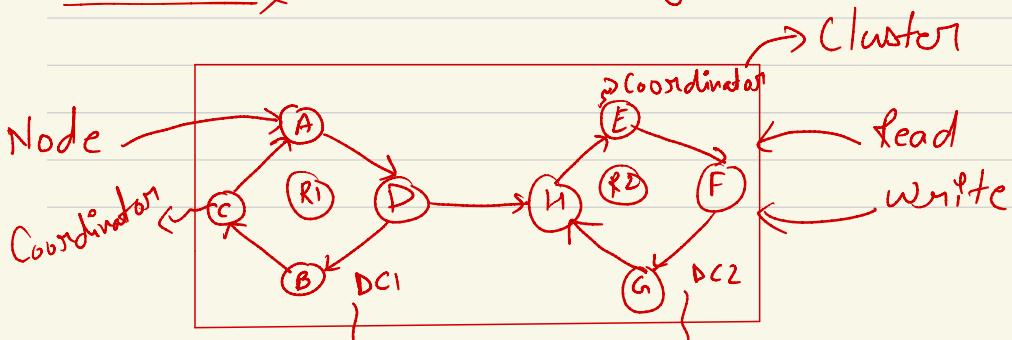
HADOOP

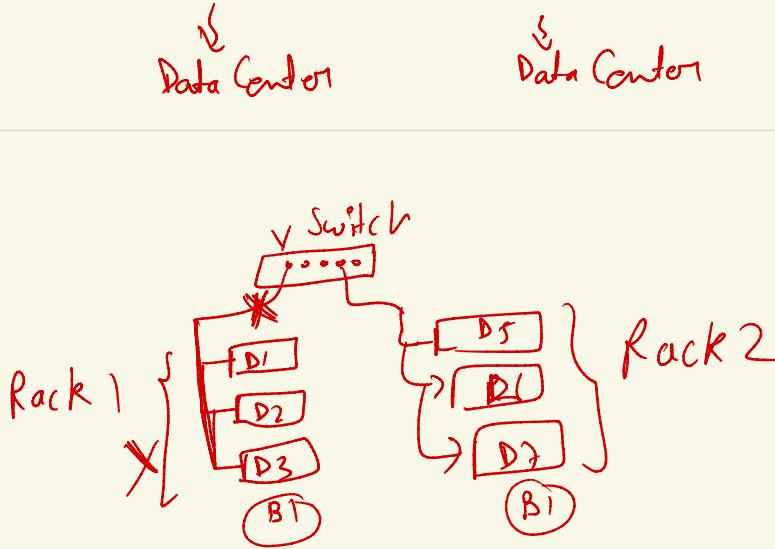


Datalake

- ↳ Means Storing Raw Data (structured, un-structured, semi-s)
- ↳ HDFS, AWS S3, Azure Blob, GCloud storage

Cassandra Architecture (Ring)





Rack Awareness

* Node → Where data stored

Cluster → Combination of Nodes

Data Center → Large location to contain Nodes

Coordinator → Manages current Read/Write request

Commit Log → Every write operation is written to Commit Log. Commit log is used for crash recovery.

Mem-Table → After data written in Commit log, data will be written in Mem-Table. Data is written in Mem-Table temporarily.

SSTable → When mem-table reaches a certain threshold, data is flushed to an SSTable disk.

Data Replication in Cassandra

↳ Means making multiple copy of Data and store it on different nodes for fault tolerance.

Cassandra places replicas of data on different nodes based on two factors.

- ① Where to place next replica is determined by replication strategy.
- ② While the total number of replicas placed on different nodes is determined by replication factor.

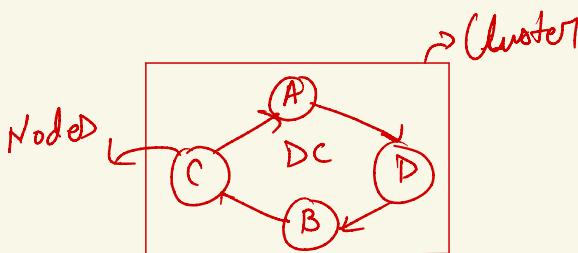
Replication Factor = How many copies of same data will be created?

Two type of Replication Strategy

- ① Simple Strategy
- ② Network Topology Strategy

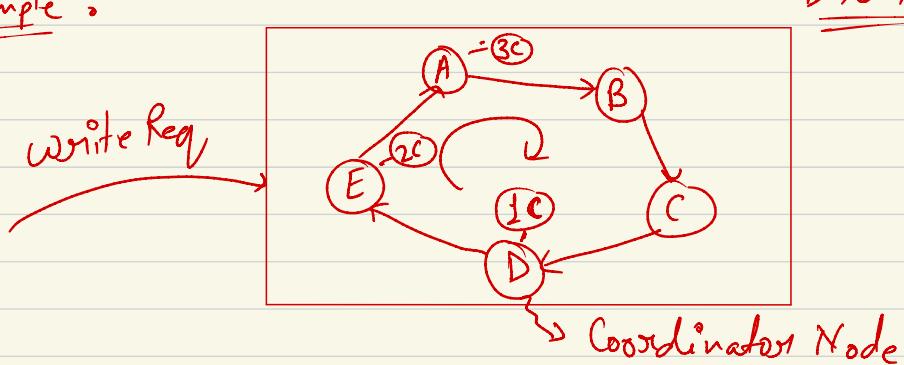
* Simple Strategy

↳ it will be used when we will one Data Center.



↳ first copy will be created on the coordinator Node and after that remaining copies will be created in other nodes in clock-wise direction.

Example : Replication factor = 3 , it will be done on D → E → A

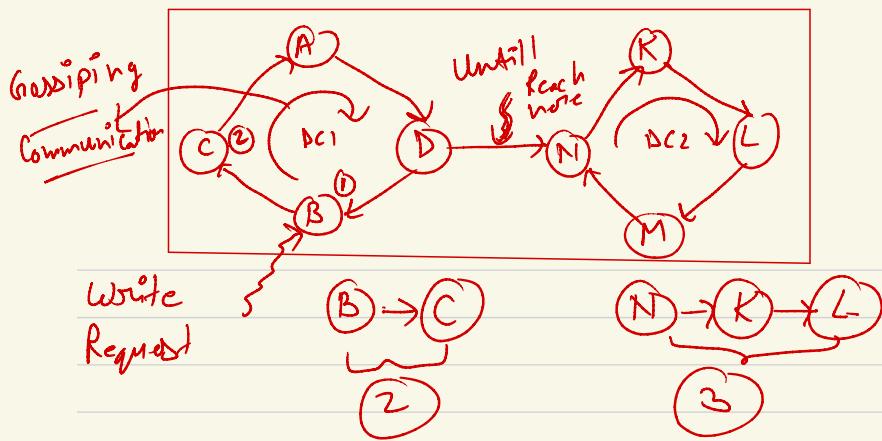


* Network Topology Strategy

↳ It is used when we have more than one data centers.

- ↳ Replicas are set for each data center separately
- ↳ Network Topology places replicas in the clockwise direction in the ring until it reaches the first node in another Data Center.

Replication Factor for DC1 = 2, for DC2 = 3



② Concept of Keys in Cassandra

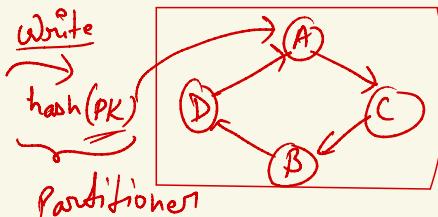
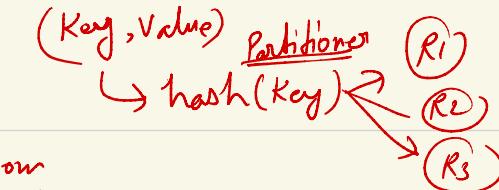
- ↳ Primary Key
- ↳ Partition Key
- ↳ Composite Partition Key
- ↳ Clustering Key

* Partition Key:

- A partition key is for

Data placement apart from
uniquely identify a record.

- Partition key will
decide on which
node data will be
stored.



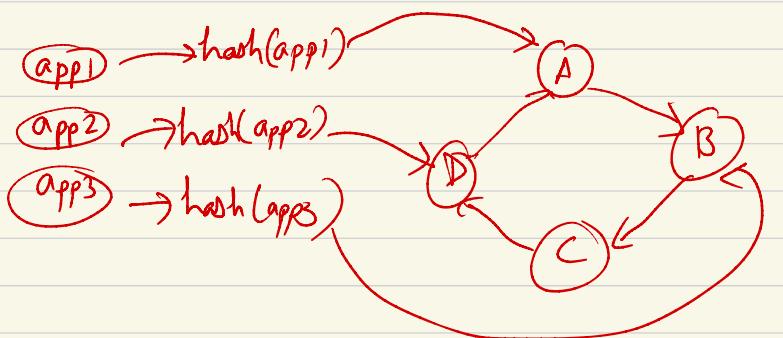
Ex: Create Table application_logs

```
( id int,  
  app_name varchar,  
  env varchar,  
  host_name varchar,  
  log_datetime timestamp,  
  log_message text,
```

PRIMARY KEY (app_name)

);

PartitionKey



② Composite Partition Key

↳ If we need to combine more than one column value to form a single partition key that will be known as Composite Partition Key.

exam Create table application_logs

(
 id,
 app-name,
 host-name,
 env,
 log-message,

PRIMARY KEY ((app-name, env))

Partition Key

* Clustering Key

↳ The columns which will be used to order the data inside a partition.

Eg Create table app-logs

(
 id,
 app-name,
 host-name,
 log-datetime,
 env,

log-message,

PRIMARY KEY (app-name, env), host-name, log-datetime)

);

Partition

Clustering Key

OR if appname is only Partition Key

PRIMARY KEY (app-name), host-name, log-datetime)

Partition
Key

Clustering Key

Partition

app-name, env, log-message, host-name, log-delete-time

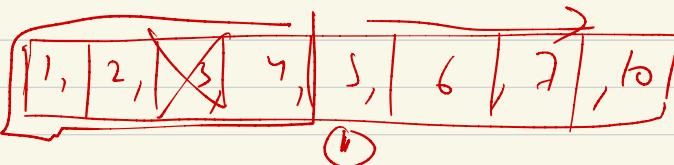
{ app1, Prod, app1INFO, host1, 2022-11-20 16:04 }

{ app1, Prod, app1INFO2, host2, 2022-11-20 16:20 }

↗? Why id is helpful?

so far search or filter related queries it can optimize.

Key = 7



↗ how to change the order for clustering?

Create table app-logs

(id,

app-name,

host-name,

env,

log-datetime,

log-message,

Primary Key ((app-name, env), host-name, log-datetime)

) with Clustering Order By (host-name Asc, log-datetime DESC),

Primary Key

↳ It consists of one or more partition key
and zero or more clustering key.

